

# DISTRIBUTED SYSTEMS ASSIGNMENT REPORT

---



## ASSIGNMENT REPORT

**Assignment ID: Assignment4 - Exploring Kubernetes**

**Student Name: 王谦益**

**Student ID: 12111003**

## DESIGN

Task 0: K8s Deployment & Service

### 1. modify the root API

1. create `app.py` and form root API, which returns pod name, pod IP, node name

```
@app.route('/')
def hello():
    pod_name = os.getenv('POD_NAME', 'unknown')
    pod_ip = os.getenv('POD_IP', 'unknown')
    node_name = os.getenv('NODE_NAME', 'unknown')
    return f"Hello from {pod_name}, IP: {pod_ip}, Node: {node_name}\n"
```

2. initial `Dockerfile` and `requirements.txt`

```
FROM python:3-slim
WORKDIR /app
COPY requirements.txt app.py ./
RUN pip3 install --no-cache-dir -r requirements.txt
CMD [ "python", "app.py" ]
```

3. write `kind-config.yaml` file to create a 4-node k8s cluster (1 control plane + 3 worker nodes)

```
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
```

```
name: flask-cluster
nodes:
- role: control-plane
- role: worker
- role: worker
- role: worker
```

4. use `t0.yaml` to create a deployment and service
5. run `docker build -t flask-app:1.0.0 .` to get the image
6. run `kind create cluster --name t0 --config kind-config.yaml` to create the cluster
7. run `kind load docker-image flask-app:1.0.0 --name t0` to load the image into the cluster
8. run `kubectl apply -f t0.yaml` to create the deployment and service
9. run `kubectl get svc` to get `ip` address and run `docker exec -it t0-control-plane /bin/bash` to start the container
10. run `curl <ip>` to get the result

## 2. graceful shutdown

```
def graceful_shutdown(signum, frame):
    print("Gracefully shutting down...")
    sys.exit(0)

signal.signal(signal.SIGTERM, graceful_shutdown)
```

## 3. greet-with-info API

1. add API in `app.py`

```
@app.route('/chat/<username>')
def greet_with_info(username):
    return f"Hello {username}, welcome to the Flask app!\n"
```

2. fix `t0.yaml` by changing the `image: flask-app:1.0.0` to `image: flask-app:1.0.1`
3. run `docker build -t flask-app:1.0.1 .` to get the new image
4. run `kind load docker-image flask-app:1.0.1 --name t0` to update the image into the cluster
5. run `kubectl apply -f t0.yaml` to change the deployment and service

6. run `kubectl get svc` to get ip address and run `docker exec -it t0-control-plane /bin/bash` to start the container

7. run `curl <ip>` and `curl <ip>/chat/<username>` to get the result

## Task 1: K8s Pod Scheduling

1. delete the deployment and service of t0
2. write `kind-config.yaml` file to create a 6-node k8s cluster (1 control plane + 5 worker nodes)

```
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
- role: control-plane
- role: worker
  labels:
    usage: normal
- role: worker
  labels:
    usage: normal
    capability: powerful
- role: worker
  kubeadmConfigPatches:
  - |
    kind: JoinConfiguration
    nodeRegistration:
    kubeletExtraArgs:
      node-labels: "usage=normal,capability=powerful"
    taints:
    - key: class
      value: "vip"
      effect: NoSchedule
- role: worker
  labels:
    usage: backup
- role: worker
  labels:
    usage: backup
```

3. use `t1.yaml` to create a deployment and service
4. run `kind create cluster --name t1 --config kind-config.yaml` to create the cluster
5. run `kind load docker-image flask-app:1.0.1 --name t1` to load the image into the cluster
6. run `kubectl apply -f t1.yaml` to create the deployment and service
7. run `kubectl scale deployment flask-app --replicas=<i>` from 1 to 5 and run `kubectl get pods -o wide` each time to get the result
8. improve `t1.yaml` by adding code

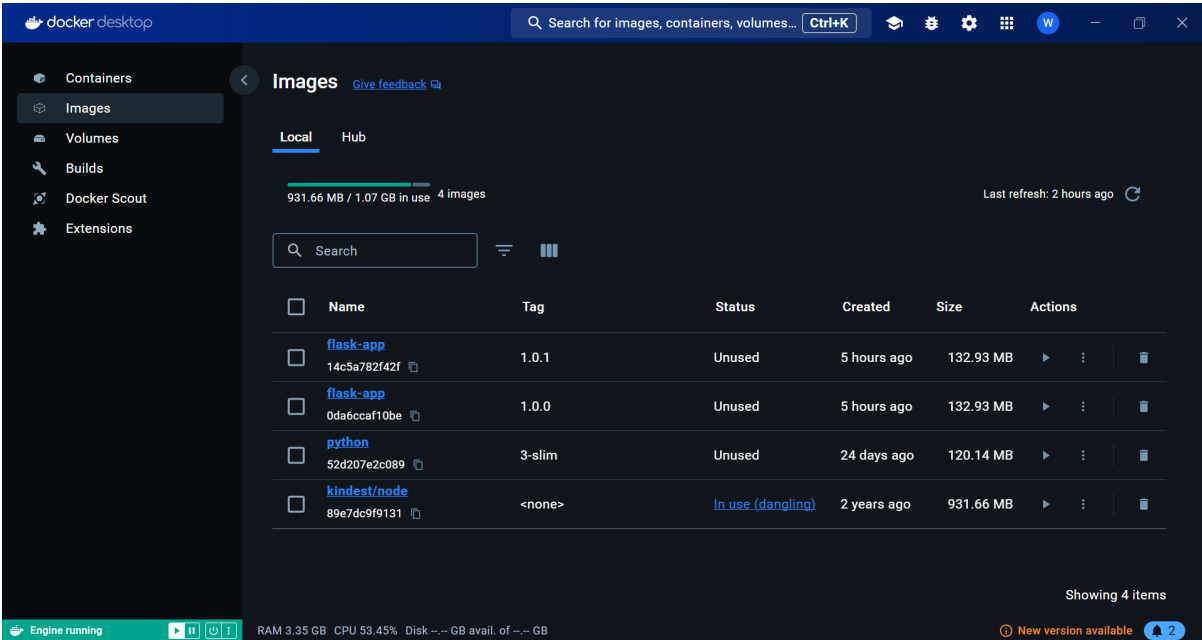
```
tolerations:
- key: "class"
  operator: "Equal"
  value: "vip"
  effect: "NoSchedule"
```

9. run `kubectl apply -f t1.yaml` to update the deployment and service and run `kubectl get pods -o wide` to get the result

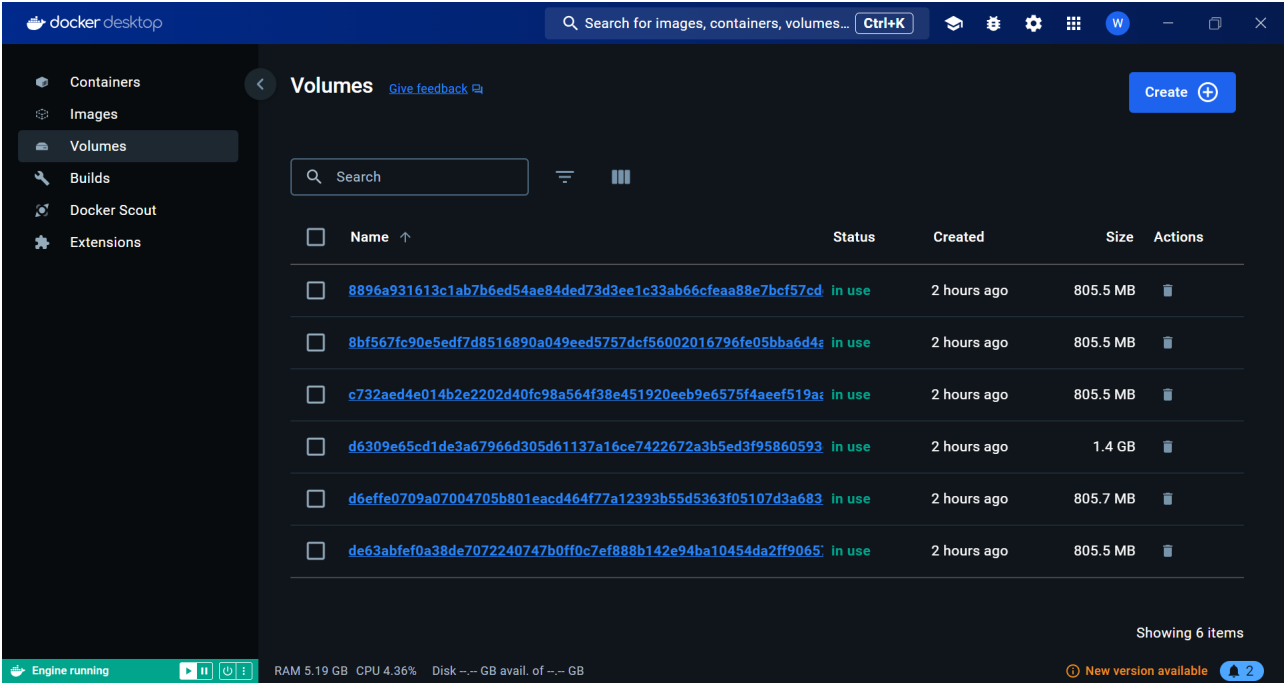
RUNNING RESULT

docker

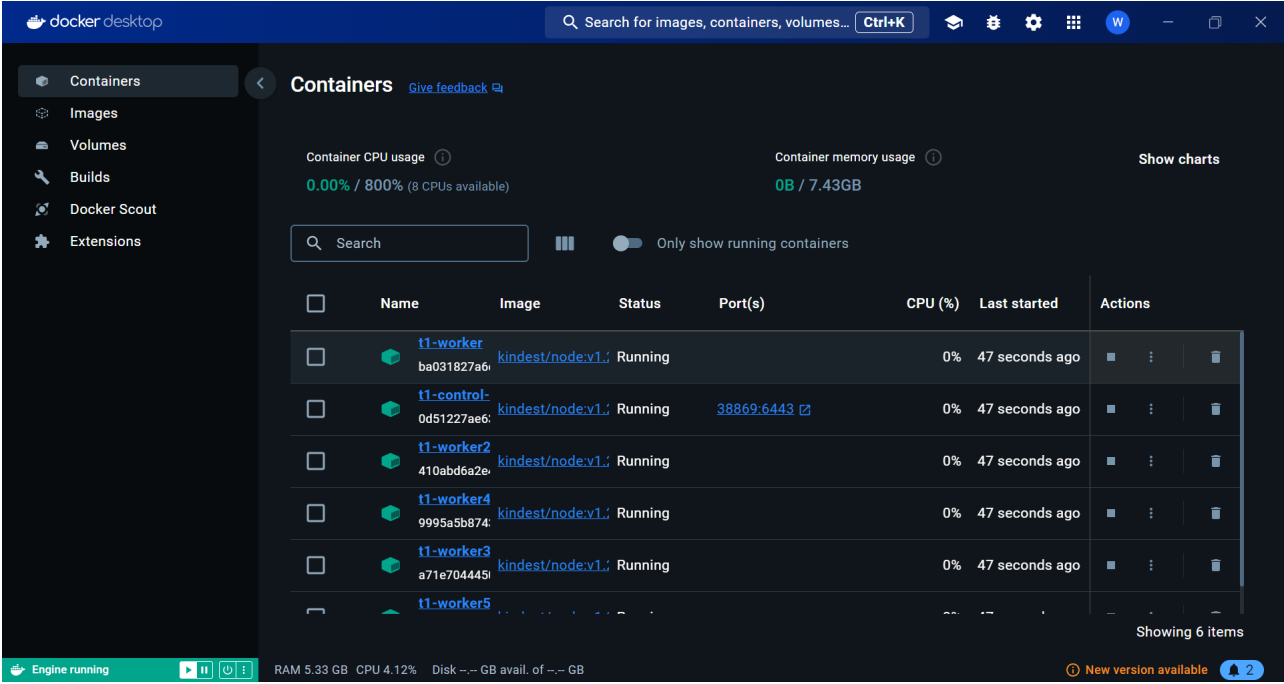
- 1. docker images
  - flask-app:1.0.0: only root API
  - flask-app:1.0.1: both root and greet-with-info API



2. docker volumes



3. docker containers(only t1 is running because t0 is already deleted)



t0

## 1. get into t0 folder, create cluster t0, load image 1.0.0, apply t0.yaml

```

(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc$ cd codebase/t0
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ kind create cluster --name t0 --config kind-config.yaml
Creating cluster "t0" ...
  ✓ Ensuring node image (kindest/node:v1.27.3)
  ✓ Preparing nodes
  ✓ Writing configuration
  ✓ Starting control-plane
  ✓ Installing CNI
  ✓ Installing StorageClass
  ✓ Joining worker nodes
Set kubectl context to "kind-t0"
You can now use your cluster with:

kubectl cluster-info --context kind-t0

Have a question, bug, or feature request? Let us know! https://kind.sigs.k8s.io/#community
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ kubectl config use-context kind-t0
Switched to context "kind-t0".
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ kind load docker-image flask-app:1.0.0 --name t0
Image: "flask-app:1.0.0" with ID "sha256:0da6ccaf10bea5c25c2b26122e4a66bc07a73fbcf861ce3091e0dd1a41d83c20" not yet present on node "t0-worker", loading...
Image: "flask-app:1.0.0" with ID "sha256:0da6ccaf10bea5c25c2b26122e4a66bc07a73fbcf861ce3091e0dd1a41d83c20" not yet present on node "t0-worker2", loading...
Image: "flask-app:1.0.0" with ID "sha256:0da6ccaf10bea5c25c2b26122e4a66bc07a73fbcf861ce3091e0dd1a41d83c20" not yet present on node "t0-worker3", loading...
Image: "flask-app:1.0.0" with ID "sha256:0da6ccaf10bea5c25c2b26122e4a66bc07a73fbcf861ce3091e0dd1a41d83c20" not yet present on node "t0-control-plane", loading...
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ kubectl apply -f t0.yaml
deployment.apps/flask-app created
service/flask-service-t0 created
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ kubectl get pods -o wide

```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
flask-app-5c685644f4-cpngh	1/1	Running	0	5s	10.244.3.2	t0-worker2	<none>	<none>
flask-app-5c685644f4-f8fc4	1/1	Running	0	5s	10.244.2.2	t0-worker	<none>	<none>
flask-app-5c685644f4-g2zvc	1/1	Running	0	5s	10.244.2.3	t0-worker	<none>	<none>
flask-app-5c685644f4-lsj89	1/1	Running	0	5s	10.244.1.2	t0-worker3	<none>	<none>

2. use `kubectl get svc` to get the ip address, start container, and curl the ip address, load new image 1.0.1, apply t0.yaml to update

```

(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ kubectl get svc

```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
flask-service-t0	ClusterIP	10.96.189.85	<none>	80/TCP	10s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	57s

```

(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ docker exec -it t0-control-plane /bin/bash
root@t0-control-plane:/# curl 10.96.189.85
Hello from flask-app-5c685644f4-cpngh, IP: 10.244.3.2, Node: t0-worker2
root@t0-control-plane:/# curl 10.96.189.85
Hello from flask-app-5c685644f4-g2zvc, IP: 10.244.2.3, Node: t0-worker
root@t0-control-plane:/# curl 10.96.189.85
Hello from flask-app-5c685644f4-lsj89, IP: 10.244.1.2, Node: t0-worker3
root@t0-control-plane:/#
exit
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ kind load docker-image flask-app:1.0.1 --name t0
Image: "flask-app:1.0.1" with ID "sha256:14c5a782f42f4a485e41a05a990983420b54e50653edec5f4de7b0f758df5bc9" not yet present on node "t0-worker", loading...
Image: "flask-app:1.0.1" with ID "sha256:14c5a782f42f4a485e41a05a990983420b54e50653edec5f4de7b0f758df5bc9" not yet present on node "t0-worker2", loading...
Image: "flask-app:1.0.1" with ID "sha256:14c5a782f42f4a485e41a05a990983420b54e50653edec5f4de7b0f758df5bc9" not yet present on node "t0-worker3", loading...
Image: "flask-app:1.0.1" with ID "sha256:14c5a782f42f4a485e41a05a990983420b54e50653edec5f4de7b0f758df5bc9" not yet present on node "t0-control-plane", loading...
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ kubectl apply -f t0.yaml
deployment.apps/flask-app configured
service/flask-service-t0 unchanged
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ kubectl get pods -o wide

```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
flask-app-5c685644f4-cpngh	1/1	Terminating	0	2m21s	10.244.3.2	t0-worker2	<none>	<none>
flask-app-5c685644f4-f8fc4	1/1	Terminating	0	2m21s	10.244.2.2	t0-worker	<none>	<none>
flask-app-5c685644f4-g2zvc	1/1	Terminating	0	2m21s	10.244.2.3	t0-worker	<none>	<none>
flask-app-5c685644f4-lsj89	1/1	Terminating	0	2m21s	10.244.1.2	t0-worker3	<none>	<none>
flask-app-845c8df6bf-cfmjv	1/1	Running	0	9s	10.244.3.3	t0-worker2	<none>	<none>
flask-app-845c8df6bf-pgbz2	1/1	Running	0	8s	10.244.1.4	t0-worker3	<none>	<none>
flask-app-845c8df6bf-s446b	1/1	Running	0	9s	10.244.2.4	t0-worker	<none>	<none>
flask-app-845c8df6bf-smzm2	1/1	Running	0	9s	10.244.1.3	t0-worker3	<none>	<none>

## 3. search ip, start container and curl the ip address again

```

(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ kubectl get svc

```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
flask-service-t0	ClusterIP	10.96.189.85	<none>	80/TCP	2m33s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3m20s

```

(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t0$ docker exec -it t0-control-plane /bin/bash
root@t0-control-plane:/# curl 10.96.189.85
Hello from flask-app-845c8df6bf-s446b, IP: 10.244.2.4, Node: t0-worker
root@t0-control-plane:/# curl 10.96.189.85
Hello from flask-app-845c8df6bf-smzm2, IP: 10.244.1.3, Node: t0-worker3
root@t0-control-plane:/# curl 10.96.189.85
Hello from flask-app-845c8df6bf-s446b, IP: 10.244.2.4, Node: t0-worker
root@t0-control-plane:/# curl 10.96.189.85/chat/a
Hello a, welcome to the Flask app!
root@t0-control-plane:/# curl 10.96.189.85/chat/b
Hello b, welcome to the Flask app!
root@t0-control-plane:/# curl 10.96.189.85/chat/c
Hello c, welcome to the Flask app!
root@t0-control-plane:/#
exit

```

## 1. get into t1 folder, create cluster t1, load image 1.0.1, apply t0.yaml

```
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc$ cd codebase/t1
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kind create cluster --name t1 --config kind-config.yaml
Creating cluster "t1" ...
  ✓ Ensuring node image (kindest/node:v1.27.3)
  ✓ Preparing nodes
  ✓ Writing configuration
  ✓ Starting control-plane
  ✓ Installing CNI
  ✓ Installing StorageClass
  ✓ Joining worker nodes
Set kubectl context to "kind-t1"
You can now use your cluster with:

kubectl cluster-info --context kind-t1

Have a question, bug, or feature request? Let us know! https://kind.sigs.k8s.io/#community
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl config use-context kind-t1
Switched to context "kind-t1".
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kind load docker-image flask-app:1.0.1 --name t1
Image: "flask-app:1.0.1" with ID "sha256:14c5a782f42f4a485e41a05a990983420b54e50653edec5f4de7b0f758df5bc9" not yet present on node "t1-worker", loading...
Image: "flask-app:1.0.1" with ID "sha256:14c5a782f42f4a485e41a05a990983420b54e50653edec5f4de7b0f758df5bc9" not yet present on node "t1-control-plane", loading...
Image: "flask-app:1.0.1" with ID "sha256:14c5a782f42f4a485e41a05a990983420b54e50653edec5f4de7b0f758df5bc9" not yet present on node "t1-worker2", loading...
Image: "flask-app:1.0.1" with ID "sha256:14c5a782f42f4a485e41a05a990983420b54e50653edec5f4de7b0f758df5bc9" not yet present on node "t1-worker4", loading...
Image: "flask-app:1.0.1" with ID "sha256:14c5a782f42f4a485e41a05a990983420b54e50653edec5f4de7b0f758df5bc9" not yet present on node "t1-worker3", loading...
Image: "flask-app:1.0.1" with ID "sha256:14c5a782f42f4a485e41a05a990983420b54e50653edec5f4de7b0f758df5bc9" not yet present on node "t1-worker5", loading...
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl apply -f t1.yaml
deployment.apps/flask-app created
service/flask-service-t1 created
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
flask-app-7c77b965dc-kcw7x	1/1	Running	0	3s	10.244.1.2	t1-worker2	<none>	<none>

## 2. scale deployment from 1 to 3 and get pods each time

```
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl scale deployment flask-app --replicas=1
deployment.apps/flask-app scaled
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
flask-app-7c77b965dc-72rlg	1/1	Terminating	0	5m53s	10.244.2.2	t1-worker	<none>	<none>
flask-app-7c77b965dc-kcw7x	1/1	Running	0	7m59s	10.244.1.2	t1-worker2	<none>	<none>
flask-app-7c77b965dc-kts4s	1/1	Terminating	0	6m3s	10.244.3.2	t1-worker5	<none>	<none>
flask-app-7c77b965dc-mb7tv	1/1	Terminating	0	5m57s	10.244.5.2	t1-worker4	<none>	<none>

```
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl scale deployment flask-app --replicas=2
deployment.apps/flask-app scaled
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
flask-app-7c77b965dc-kcw7x	1/1	Running	0	8m44s	10.244.1.2	t1-worker2	<none>	<none>
flask-app-7c77b965dc-r94gp	1/1	Running	0	2s	10.244.3.3	t1-worker5	<none>	<none>

```
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl scale deployment flask-app --replicas=3
deployment.apps/flask-app scaled
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
flask-app-7c77b965dc-kcw7x	1/1	Running	0	8m47s	10.244.1.2	t1-worker2	<none>	<none>
flask-app-7c77b965dc-r94gp	1/1	Running	0	5s	10.244.3.3	t1-worker5	<none>	<none>
flask-app-7c77b965dc-w4drz	1/1	Running	0	1s	10.244.5.3	t1-worker4	<none>	<none>

## 3. scale deployment from 3 to 5 and get pods each time, find a pod is not ready

```
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl scale deployment flask-app --replicas=4
deployment.apps/flask-app scaled
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
flask-app-7c77b965dc-9th7g	1/1	Running	0	2s	10.244.2.3	t1-worker	<none>	<none>
flask-app-7c77b965dc-kcw7x	1/1	Running	0	8m51s	10.244.1.2	t1-worker2	<none>	<none>
flask-app-7c77b965dc-r94gp	1/1	Running	0	9s	10.244.3.3	t1-worker5	<none>	<none>
flask-app-7c77b965dc-w4drz	1/1	Running	0	5s	10.244.5.3	t1-worker4	<none>	<none>

```
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl scale deployment flask-app --replicas=5
deployment.apps/flask-app scaled
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
flask-app-7c77b965dc-9th7g	1/1	Running	0	6s	10.244.2.3	t1-worker	<none>	<none>
flask-app-7c77b965dc-fqg65	0/1	Pending	0	1s	<none>	<none>	<none>	<none>
flask-app-7c77b965dc-kcw7x	0/1	Running	0	8m55s	10.244.1.2	t1-worker2	<none>	<none>
flask-app-7c77b965dc-r94gp	1/1	Running	0	13s	10.244.3.3	t1-worker5	<none>	<none>
flask-app-7c77b965dc-w4drz	1/1	Running	0	9s	10.244.5.3	t1-worker4	<none>	<none>

## 4. improve t1.yaml and apply t1.yaml to update

```
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl apply -f t1.yaml
deployment.apps/flask-app configured
service/flask-service-t1 unchanged
(dcc_ass4) (base) win@DESKTOP-Q3V6AG6:~/dcc/codebase/t1$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
flask-app-7c77b965dc-9th7g	1/1	Terminating	0	20m	10.244.2.3	t1-worker	<none>	<none>
flask-app-7c77b965dc-kcw7x	1/1	Terminating	0	29m	10.244.1.2	t1-worker2	<none>	<none>
flask-app-7c77b965dc-r94gp	1/1	Terminating	0	20m	10.244.3.3	t1-worker5	<none>	<none>
flask-app-7c77b965dc-w4drz	1/1	Terminating	0	20m	10.244.5.3	t1-worker4	<none>	<none>
flask-app-847998c9c9-xvzqx	1/1	Running	0	7s	10.244.4.2	t1-worker3	<none>	<none>

## PROBLEMS

when write `kind-config.yaml` for t1 in the first time, I use:

```
- role: worker
  labels:
    usage: normal
```

```
    capability: powerful
  taints:
    - key: class
      value: "vip"
      effect: NoSchedule
```

but there's error:

```
ERROR: failed to create cluster: unable to decode config: yaml: unmarshal errors
```

then I check the slides and change the code to:

```
- role: worker
  kubeadmConfigPatches:
  - |
    kind: JoinConfiguration
    nodeRegistration:
      kubeletExtraArgs:
        node-labels: "usage=normal,capability=powerful"
    taints:
      - key: class
        value: "vip"
        effect: NoSchedule
```

then it works

## Advice on Future Cloud Computing Lab

As for me, I think the lab is good, but I think it can be improved by adding more details and examples about cloud platform or cloud security measures and data storage solutions maybe.