

A decorative graphic on the left side of the slide, consisting of a network of light blue lines and small circles, resembling a circuit board or a neural network, extending from the top to the bottom.

DISTRIBUTED AND CLOUD COMPUTING

LAB1: MPI INTRODUCTION AND SETUP

GEORGE DIAMANTOPOULOS

OVERVIEW

- During the lab:
 - Introduction to technical concepts in Distributed and Cloud Computing
 - Application of concepts and Q&A
- Programming language: C
- Blackboard site:
 - [Distributed and Cloud Computing Fall 2024](https://bb.sustech.edu.cn) : **bb.sustech.edu.cn**

DISTRIBUTED COMPUTING INTUITION

Doing an assignment alone VS doing an assignment in a group?

Can you think of one **good** thing and one **bad** thing for each?

DISTRIBUTED COMPUTING INTUITION

Alone:

- + Management is simple
- You have to do everything

Group:

- + Work is split amongst many people (ideally!)
- You have to communicate and assign tasks

DISTRIBUTED COMPUTING INTUITION

Same things apply to computers!

MPI – MESSAGE PASSING INTERFACE



- Core requirement for all parallel systems: **Communication**
 - **Shared Memory**
 - **Message Passing**
- **MPI**: a message passing **standard** to allow for distributed/parallel computation
 - Many implementations [open-mpi, mpich, ...] for various languages [C, C++, FORTRAN, ...]
- This course: **open-MPI** for **C**
- MPI is **extremely common** for High Performance Computing (HPC) applications
 - If something runs on supercomputer (and does not use GPUs) there's a good chance it uses MPI

MPI – MESSAGE PASSING INTERFACE



- MPI implements an interface for *parallel process communication*
 - Abstracts the low-level details of processes communication
 - Allows the programmer to focus on the problem at hand (the parallel application)!
- You can think of **MPI** as a **postal service** and the processes as people sending letters
 - When you want to send a letter, you just drop it in the mailbox!


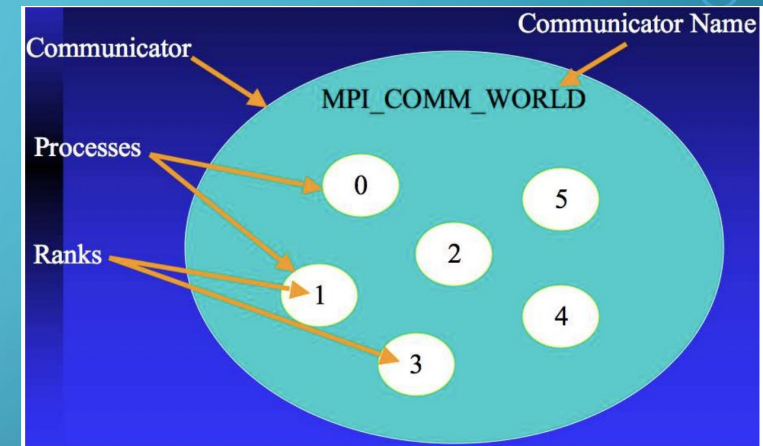
MPI - PROCESSES



- **MPI processes are managed by MPI and run concurrently (at the same time)**
 - On the CPU cores of a computer or server
 - On the cores of many CPUs each in a cluster of network computers
- **Each process:**
 - Runs an instance of the parallel application (e.g. large weather model)
 - Has its own local memory (does not share memory with any other process)
 - Can use the MPI interface to send messages to other processes

MPI - PROCESSES

- MPI communicators - (neighbourhoods)
 - Group processes
 - Assign them unique ID's called **ranks** (house addresses)
- "MPI_COMM_WORLD" is the default communicator
 - Contains all MPI processes generated for the application
- Custom communicators can be defined by the user
 - Can contain just a subset of the processes
 - Use full for managing the complex communication patterns



SETTING UP THE ENVIRONMENT FOR MPI

WINDOWS

• **Opt. 1: WSL2 (RECOMMENDED)**

- Open the windows store and search 'Ubuntu' and install
- Open ubuntu app and set up username and password when prompted

• **Opt. 2: Ubuntu Virtual Machine**

- Download a VM hypervisor
 - VirtualBox is free and recommended
- Download an ISO of the latest ubuntu LTS
- There are plenty of guides online to help you

MAC OS

• **Opt. 1: Install locally (RECOMMENDED)**

- Through homebrew or building from source

• **Opt. 2: Ubuntu Virtual Machine**

- **not recommended:** macs natively support openMPI
- Similar to Windows
- Intel mac: VirtualBox
- Apple silicone: VMware fusion



Canonical Group Limited

SETTING UP THE ENVIRONMENT FOR MPI

Ubuntu install (WLS or VM)

Run the following commands in the terminal:

```
sudo apt-get update
```

```
sudo apt install gcc (C compiler: if not installed)
```

```
sudo apt-get install openmpi-bin openmpi-doc libopenmpi-dev
```

MAC OS Install

Get xcode-command line tools: `xcode-select --install`

Opt. 1: Homebrew (RECOMMENDED)

- Install homebrew packet manager: <https://brew.sh/>
- Run: `brew install open-mpi`

If download is slow or blocked set up the brew sustech mirror:

<https://mirrors.sustech.edu.cn/help/homebrew.html#introduction>

Opt. 2: Build openMPI (NOT RECOMMENDED)

- First answer at: <https://www.open-mpi.org/faq/?category=building>
- Download the zipped source code and unzip
- Go the unzipped directory and run the configuration
- Make all install

Code example:

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char* argv[]) {
    // Initialization
    MPI_Init(NULL, NULL);
    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);
    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d processors\n",
           processor_name, world_rank, world_size);
    // Finalize the MPI environment.
    MPI_Finalize();
}
```

MPI - BOILERPLATE CODE

```
#include <mpi.h>
// other libraries and header files

int main(int argc, char* argv[]) {
    // Initialization
    MPI_Init(NULL, NULL);

    // PARALLEL APPLICATION LOGIC.

    // Finalize the MPI environment.
    MPI_Finalize();
}
```

Imports the MPI library

Initialises the MPI environment

Finalise the MPI environment

THIS CODE NEVER CHANGES AND IS ALWAYS REQUIRED WHEN WORKING WITH MPI

MPI – USEFUL UTILITY FUNCTIONS

```
// Get the number of processes
int world_size;
MPI_Comm_size(MPI_COMM_WORLD, &world_size);
```

Sets world_size equal to the total number of MPI processes

```
// Get the rank of the process
int world_rank;
MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
```

Sets world_rank equal to the rank of the current process

```
// Get the name of the processor
char processor_name[MPI_MAX_PROCESSOR_NAME];
int name_len;
MPI_Get_processor_name(processor_name,
&name_len);
```

Retrieves the name of the processor (usually, the name you have given to the PC like GeorgesLaptop)

Code example:

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char* argv[]) {
    // Initialization
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

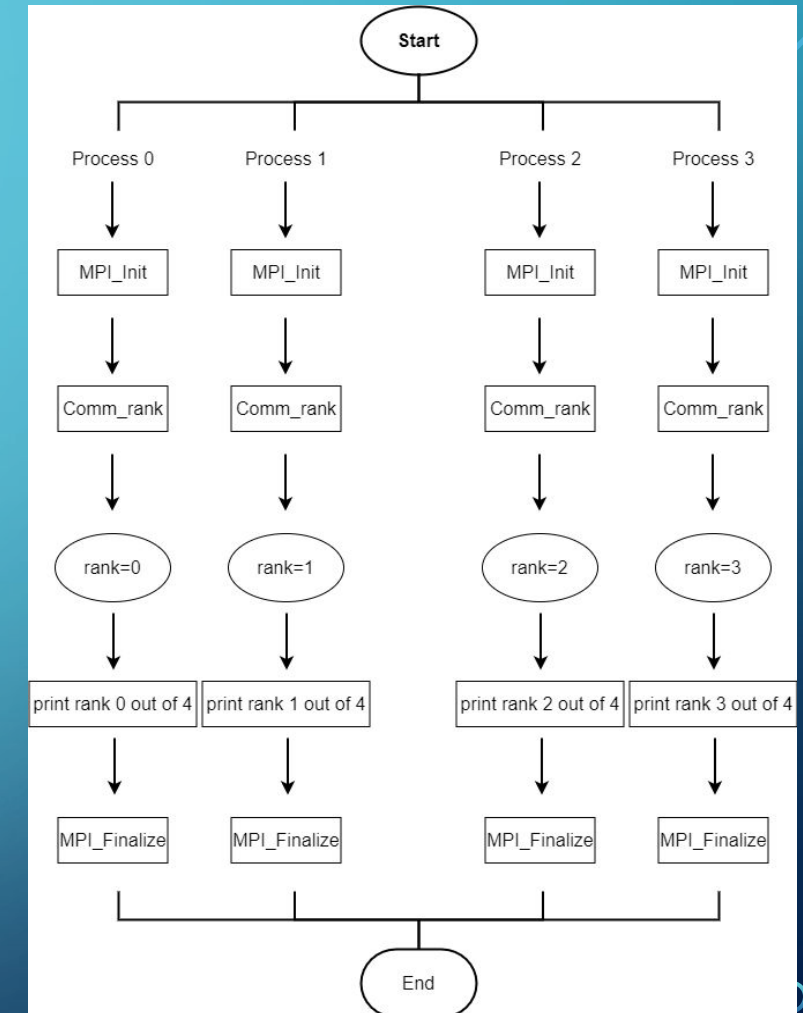
    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d\n",
           processor_name, world_rank, world_size);

    // Finalize the MPI environment.
    MPI_Finalize();
}
```

Diagram of MPI execution



https://github.com/wesleykendall/mpitutorial/blob/g_h-pages/tutorials/mpi-hello-world/code/mpi_hello_world.c

COMPILATION AND RUNNING

To compile, use **mpicc**, which is just a wrapper around gcc:

```
mpicc source.c -o executable_name
```

Then run the program use **mpirun**:

```
mpirun -np 4 ./executable_name
```

-np X: tells MPI how many processes to create

The result:

```
george@DESKTOP-E24BUDU:~/DistSys24/MPI/w1$ mpicc mpi_intro.c -o intro
george@DESKTOP-E24BUDU:~/DistSys24/MPI/w1$ mpirun -np 8 ./intro
Hello world from processor DESKTOP-E24BUDU, rank 7 out of 8 processors
Hello world from processor DESKTOP-E24BUDU, rank 0 out of 8 processors
Hello world from processor DESKTOP-E24BUDU, rank 4 out of 8 processors
Hello world from processor DESKTOP-E24BUDU, rank 5 out of 8 processors
Hello world from processor DESKTOP-E24BUDU, rank 6 out of 8 processors
Hello world from processor DESKTOP-E24BUDU, rank 1 out of 8 processors
Hello world from processor DESKTOP-E24BUDU, rank 2 out of 8 processors
Hello world from processor DESKTOP-E24BUDU, rank 3 out of 8 processors
george@DESKTOP-E24BUDU:~/DistSys24/MPI/w1$
```

MPI PROGRAMMING INTUITION

- Every process runs the same program! What's the point?
 - **Ranks!**
- Every process has a unique rank which we can use to assign it different tasks!

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char* argv[]) {
    // Initialization
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    if(world_rank % 2 == 0){
        // Print a hello world message
        printf("Hello world from processor %s, rank %d out of %d\n", processor_name, world_rank, world_size);
    }

    // Finalize the MPI environment.
    MPI_Finalize();
}
```



```
george@DESKTOP-E24BUDU:~/DistSys24/MPI/w1$ mpicc mpi_intro_plus.c -o plus
george@DESKTOP-E24BUDU:~/DistSys24/MPI/w1$ mpirun -np 8 ./plus
Hello world from processor DESKTOP-E24BUDU, rank 2 out of 8 processors
Hello world from processor DESKTOP-E24BUDU, rank 4 out of 8 processors
Hello world from processor DESKTOP-E24BUDU, rank 6 out of 8 processors
Hello world from processor DESKTOP-E24BUDU, rank 0 out of 8 processors
george@DESKTOP-E24BUDU:~/DistSys24/MPI/w1$
```

COMPILATION AND RUNNING (EXTRA)

- **Process to CPU core mapping**

- MPI runs new processes in what are call 'slots'
- By default, MPI identifies the number of CPU cores and creates an equal number of **slots**
- In this case, due to multithreading, there are 16 threads available for processing
- Going over that amount is possible, but processes are not fully parallel (CPU task scheduling)

```
george@DESKTOP-E24BUDU:~/DistSys24/MPI/w1$ mpirun -np 16 ./intro
```

```
-----  
There are not enough slots available in the system to satisfy the 16  
slots that were requested by the application:
```

```
./intro
```

```
Either request fewer slots for your application, or make more slots  
available for use.
```

```
A "slot" is the Open MPI term for an allocatable unit where we can  
launch a process. The number of slots available are defined by the  
environment in which Open MPI processes are run:
```

```
george@DESKTOP-E24BUDU:~/DistSys24/MPI/w1$ mpirun -np 16 --oversubscribe ./intro  
Hello world from processor DESKTOP-E24BUDU, rank 5 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 8 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 9 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 10 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 3 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 13 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 14 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 15 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 0 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 11 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 12 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 4 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 7 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 1 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 6 out of 16 processors  
Hello world from processor DESKTOP-E24BUDU, rank 2 out of 16 processors
```


CONFIGURATION OF MPIRUN

- Official documentation available at:
 - <https://www.open-mpi.org/doc/v4.0/man1/mpirun.1.php>

Structure of mpirun: `mpirun [MPI options] <program> [program arguments]`

KEY TAKEAWAY: Be careful to add MPI options **BEFORE** the program, otherwise, they will be treated as arguments for the program and will not take effect!