DISTRIBUTED AND CLOUD COMPUTING

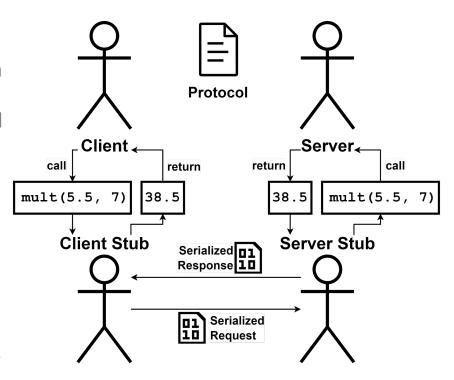
LAB 8: REVERSE PROXY + OTHER API ARCHITECTURES

(Module: RPC & RESTFUL API)

RPC Stub as A Proxy

- Client Stub
 - a. Handles serialization / deserialization of data for the client.
 - b. Handles remote server connection and data transmission for the client.
- Server Stub
 - a. Works similarly, but for the server.

- RPC Stubs act like proxies.
- The concept of a proxy is broad, but in computer networking, a **forward proxy** and a **reverse proxy** each have specific, distinct functions.

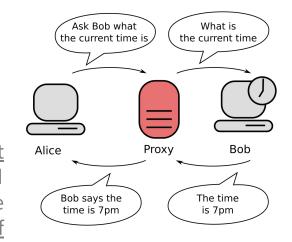


Forward Proxy

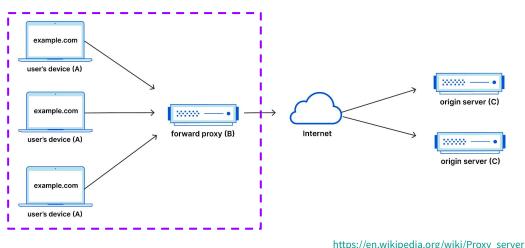
"A forward proxy, often called a proxy, proxy server, or web proxy, is a <u>server that sits in front of a group of client machines</u>. When those computers make requests to sites and services on the Internet, the proxy server intercepts those requests and then communicates with web servers <u>on behalf of those clients</u>, like a middleman."



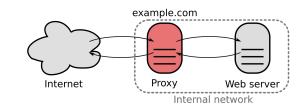
- 1. Access Control & Security
- 2. Privacy & Anonymity
- 3. Content Caching & Filtering
- 4. Monitoring & Logging
- 5. Transparency



Forward Proxy Flow



Reverse Proxy

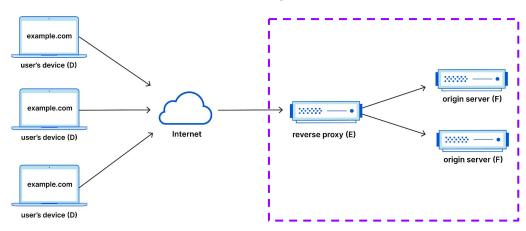


- "A reverse proxy is a <u>server that sits in front of one or more web servers</u>, intercepting requests from clients."
- It works for the <u>servers</u> compared to the forward proxy.
- Forward & Reverse Proxy are coexisting twins.
- Example: Load Balancers, CDNs, API Gateways, Web Application Firewalls (WAFs), etc.

Why Reverse Proxy? DDos attack

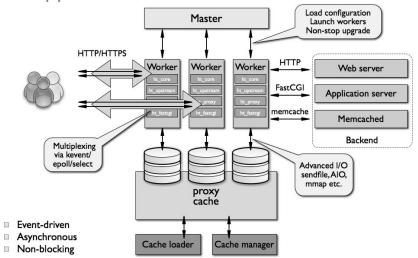
- 1. Load Balancing
- 2. Access Control & Security
- 3. Content Caching & Filtering
- 4. Monitoring & Logging
- 5. SSL Encryption Offloading

Reverse Proxy Flow



Nginx as A Reverse Proxy

- "nginx ("engine x") is an HTTP web server, reverse proxy, content cache, load balancer, TCP/UDP proxy server, and mail proxy server."
- one of the most popular tools for load balancing purpose
- simple configuration file format
- clean support with Docker





```
nginx.conf X
rpc rest > 2 others > 0 reverse proxy > nginx > 😃 nginx.conf
      # https://nginx.org/en/docs/http/load balancing.html
      # placeholder: event handlers
      events {}
      # handle HTTP requests
      http 🖟
        # define a group of backend servers that will handle requests
        upstream flask servers {
          # Default load balancing method is round-robin.
          # To test other methods, uncomment one of the following:
          # least conn; # forward to the server with the least number
                          # same ip-hash to the same server (session main
          server flask server1:5000
          server flask server2:5000
          server flask server3:5000:
        # virtual server config
        server {
          listen 80
          # process requests matching the URI pattern ('/' means root +
          location /
            # forward requests to the upstream group
            proxy pass http://flask servers
            # note back info of the original client
            # $remote addr: https://nginx.org/en/docs/http/ngx http core
            # $proxy add x forwarded for: https://nginx.org/en/docs/http
            # X-Forwarded-For records a list of hopped IPs, compared to
            proxy set header Host $host:
            proxy set header X-Real-IP $remote addr;
            proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for
            proxy set header X-Forwarded-Proto $scheme;
```

Load balance requests to 3 identical Flask servers.

- > Reference codebase: revproxy_nginx
 - 1. Set up Python (Miniconda is recommended).
- 2. Check flask_app/app.py for the Flask server implementation. It contains 1 simple API at the base URL that returns a Hello World message marking the server's id/name.
- 3. Check flask_app/Dockerfile for the Flask server image specification. Build the image:
 - docker build -t flask_app .

 $\label{times} \mbox{TIPS:} \ \mbox{\bf docker image prune} \ \mbox{to clean previous builds}.$

```
papp.py x

rpc_rest > 2_others > 0_reverse_proxy > flask_app > \( \Phi \) app.py > ...

from flask import Flask

import os

app = Flask(__name__)

@app.route('/')

def hello():

server_name = os.getenv('SERVER_NAME', 'Unknown Server')

return f'Hello from {server_name}!'

if __name__ == '__main__':

app.run(host='0.0.0.0', port=int(os.getenv('FLASK_PORT', 5000)))

app.run(host='0.0.0.0', port=int(os.getenv('FLASK_PORT', 5000)))
```

```
Dockerfile X

rpc_rest > 2_others > 0_reverse_proxy > flask_app > → Dockerfile > ...

    FROM python:3-slim

    WORKDIR /app

COPY requirements.txt app.py ./

RUN pip3 install --no-cache-dir -r requirements.txt

CMD [ "python", "app.py" ]

CMD [ "python", "app.py" ]

Proceedings of the process of
```

Load balance requests to 3 identical Flask servers.

- > Reference codebase: revproxy_nginx
 - Check nginx/nginx.conf that configures how nginx server intercepts & forwards the requests.
 - a. Listen for new requests on port 80.
 - b. Process requests matching / pattern only.
 - c. Forward requests to upstream server group.

```
nainx.conf X
rpc_rest > 2_others > 0_reverse_proxy > nginx > 🗘 nginx.conf
      # placeholder: event handlers
      events {}
      # handle HTTP requests
      http /
        # define a group of backend servers that will handle requests
        upstream flask servers {
          # Default load balancing method is round-robin.
          # To test other methods, uncomment one of the following:
          # least conn; # forward to the server with the least number
                           # same ip-hash to the same server (session main
          server flask server1:5000;
          server flask server2:5000:
          server flask server3:5000;
        # virtual server config
        server {
          listen 80:
          # process requests matching the URI pattern ('/' means root + e
          location / {
            # forward requests to the upstream group
            proxy_pass_http://flask_servers;
            # $host: https://nginx.org/en/docs/http/ngx http core module
            # $remote addr: https://nginx.org/en/docs/http/ngx http core
            # $proxy_add_x_forwarded_for: https://nginx.org/en/docs/http
            # $scheme: https://nginx.org/en/docs/http/ngx http core modu
            # X-Forwarded-For records a list of hopped IPs, compared to
            proxy set header Host $host;
            proxy_set_header X-Real-IP $remote_addr;
            proxy set header X-Forwarded-For $proxy add x forwarded for
            proxy_set_header X-Forwarded-Proto $scheme;
```

Load balance requests to 3 identical Flask servers.

- > Reference codebase: revproxy_nginx
- 4. Check nginx/nginx.conf that configures how nginx server intercepts & forwards the requests.
 - a. Listen for new requests on port 80.
 - b. Process requests matching / pattern only.
 - c. Forward requests to upstream server group.
- 5. Check compose.yaml for the defined services.
 - a. 3 Flask servers serve at port 5000, assigned with a server name.
 - b. 1 nginx server loads the nginx.conf file, serving at port 80 which is exposed to localhost.
- 6. Start the environment:
 - a. docker compose up -d

```
compose,yaml 1 X
rpc_rest > 2_others > 0_reverse_proxy > * compose.yaml > {} services > {}
       docker-compose.yml - The Compose specification establishes a standard f
         # Flask servers
         flask server1:
           image: flask_app
                                # local build
           environment:
             FLASK PORT: 5000
             SERVER NAME: "Flask Server 1"
         flask server2:
           image: flask app
           environment:
             FLASK PORT: 5000
             SERVER NAME: "Flask Server 2"
         flask_server3:
           image: flask_app
           environment:
             FLASK PORT: 5000
             SERVER_NAME: "Flask Server 3"
         # Nginx for load balancing
           image: nginx:1.27
           volumes:
             # https://docs.docker.com/reference/compose-f
             - ./nginx/nginx.conf:/etc/nginx/nginx.conf:ro
           ports:
             - "80:80"
           depends on:
                                               read-onl
             - flask server1
             - flask server2
             - flask server3
```

Load balance requests to 3 identical Flask servers.

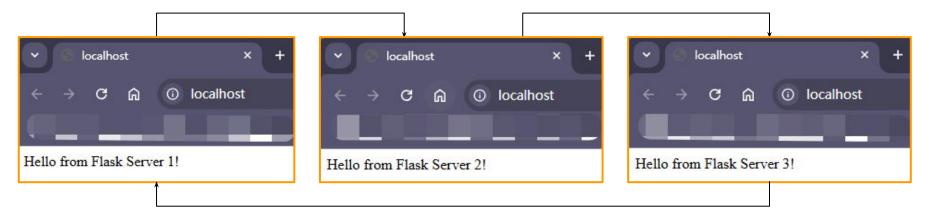
- > Reference codebase: revproxy_nginx
- 7. Start a web browser and access http://localhost:80/. Refresh the page multiple times. Note how the requests are handled by different Flask servers in a round-robin fashion (default load balancing approach of nginx).

✓Network 0_reverse_proxy_default

Container 0_reverse_proxy-flask_server3-1 Started Container 0_reverse_proxy-flask_server2-1 Started

Container 0_reverse_proxy-flask_serverl-1
Container 0_reverse_proxy-nginx-1

- 8. Optionally, test with Postman or cURL.
 - a. e.g., curl localhost



Load balance requests to 3 identical Flask servers.

- > Reference codebase: revproxy_nginx
 - 9. Check test/clients.py where 1000 requests are sent concurrently to the nginx server. Then, a map is constructed to record how many requests are handled in each Flask server. Run:
 - python test/clients.py

```
python test/clients.py
{1: 333, 2: 333, 3: 334}
```

```
python test/clients.py
0: Hello from Flask Server 1!
1: Hello from Flask Server 3!
2: Hello from Flask Server 2!
3: Hello from Flask Server 2!
8: Hello from Flask Server 3!
5: Hello from Flask Server 1!
4: Hello from Flask Server 3!
7: Hello from Flask Server 1!
6: Hello from Flask Server 2!
9: Hello from Flask Server 1!
{1: 4, 2: 3, 3: 3}
```

10 concurrent requests

```
clients.py X
rpc_rest > 2_others > 0_reverse_proxy > test > 🍖 clients.py > ...
       from concurrent.futures import ThreadPoolExecutor, as completed
       from pprint import pprint
       import requests
       def send request(id, url):
           response = requests.get(url)
           msg = response.text.strip()
           server_id = int(msg.rsplit(' ', maxsplit=1)[-1][:-1])
         except requests.RequestException as e:
           msg = str(e)
           server id = -1
         return server_id, f'{id}: {msg}'
       def test_load_balancing(url, n_requests):
         # server_id => cnt
         res dict = {}
         with ThreadPoolExecutor(max workers=n requests) as executor:
           res_futures = [executor.submit(send_request, i, url) for i in range(n_requests)
           for res_future in as_completed(res_futures):
             res server id, res msg = res future.result()
             # statistics
             if res server id not in res dict:
               res_dict[res_server_id] = 0
             res_dict[res_server_id] += 1
             # sys.stdout.write(f'{res_msg}\n')
             # sys.stdout.flush()
         pprint(res_dict)
       if __name__ == '__main__':
         test_load_balancing(url='http://localhost:80/', n_requests=1000)
```

Load balance requests to 3 identical Flask servers.

- > Reference codebase: revproxy_nginx
 - 9. Check test/clients.py where 1000 requests are sent concurrently to the nginx server. Then, a map is constructed to record how many requests are handled in each Flask server. Run:
 - o python test/clients.py
- 10. Switch to other load balancing methods by modifying nginx.conf. Restart the nginx server container manually to refresh and run the test script again.
- 11. Clean up the resources via:
 - docker compose down

```
pc_rest > 2_others > 0_reverse_proxy > nginx > 0 nginx.conf
6 http {
8     upstream flask_servers {
9     # Default load balancing method is round-robin.
10     # To test other methods, uncomment one of the following:
11     # least_conn; # forward to the server with the least number of active connections
12     # ip_hash; # same ip-hash to the same server (session maintaining)
13     server flask_server1:5000;
14     server flask_server3:5000;
15     server flask_server3:5000;
16     }
```

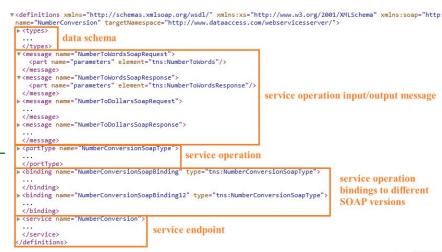
```
python test/clients.py
{1: 340, 2: 333, 3: 327}

python test/clients.py
{1: 1000}

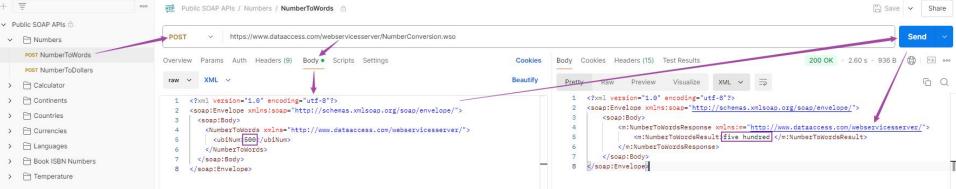
ip_hash
```

```
# docker ps
CONTAINER ID
                                                       CREATED
                                                                        STATUS
01ddcb2360bd
               nainx:1.27
                             "/docker-entrypoint..."
                                                      22 seconds ago
                                                                                        0.0.0.0:80->80/tcp
                                                                        Up 20 seconds
                                                                                                               0_reverse_proxy-nginx-1
f7f57b4f01a6
                             "python app.py"
                                                                                                               0_reverse_proxy-flask_server3-1
                             "python app.py"
                                                                                                               0_reverse_proxy-flask_server2-1
32ac3acfe6c2
                             "python app.py"
                                                                                                               0_reverse_proxy-flask_server1-1
(base)
                                                                                        # docker container restart 01ddcb2360bd
01ddcb2360bd
```

- Use XML messages + a predefined contract
- Strict, complex and verbose
- Play with a <u>Postman workspace</u> about SOAP.
- Check its WSDL definition by appending ?WSDL to the service URL, or click this link.



XML



SOAP

XML-based

for enterprise application

GraphQL

- queries exactly needed within a single request
- suitable for complex data requirements
- Explore <u>Countries GraphQL API</u>.
- Check schema definition <u>here</u>.

```
1 → query GetCountry {
      country(code: "BR") {
        native
        capital
        emoji
        currency
        languages {
          code
10
          name
12
13
Variables
          Headers
```

```
125 V
                                                                 resolve: (country) =>
                    lang: t.arg.string({
                                                                   country.languages.map((code) => ({
                      validate: (lang) => langs(
                                                                     ...languages[code as keyof typeof languages],
                    }),
                                                 128
                                                                    code,
                  resolve: async (country, { lang 129
                                                                   })),
                                                              }),
                }),
                 native: t.exposeString("native"),
                 phone: t.exposeString("phone")
                                                                                 + countries.trevorblades.com
   "data": {
     "country": {
       "name": "Brazil".
       "native": "Brasil",
       "capital": "Brasília",
       "emoji": "BR",
       "currency": "BRL",
       "languages": [
            "code": "pt",
            "name": "Portuguese"
                                                                                        RequestID: 2382755471883397900
IP limit 47 (out of 50) requests remaining (refills in 14s)
```

Payload

Headers

▼ General

Request URL:

Status Code:

Connection:

Upgrade:

▼ Request Headers

Request Method:

▼ Response Headers

Sec-Websocket-Accept:

Sec-Websocket-Extensions:

Messages

Raw

Initiator

GET

Upgrade

WebSocket

Timina

101 Switching Protocols

DEDiKVKDF1IGupnY3KaGE1OTTiU=

initial HTTP GET request

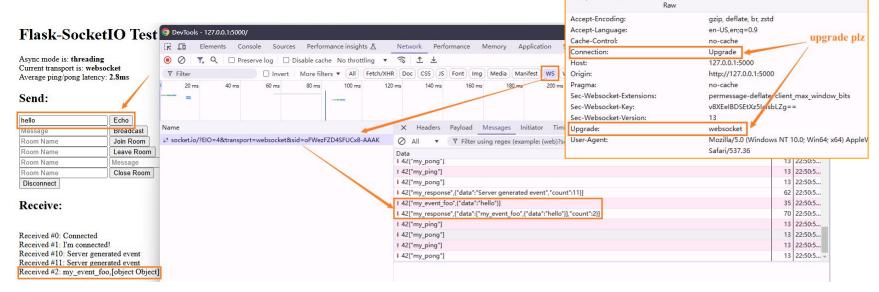
connection upgraded

ws://127.0.0.1:5000/socket.io/?EIO=4&transport=

permessage-deflate; client max window bits=15

WebSocket

- enables fast, bidirectional, and persistent connections
- benefits live chat apps and real-time gaming
- Core: upgrade initial HTTP conn to WebSocket conn.
- Try <u>Flask-SocketIO live chat demo</u>.



① Security

@ General

83 Collaborators

Code and automation

Moderation options

Access

Settings

We'll send a post request to the URL below with details of any sul

format you'd like to receive (JSON, x-www-form-urlencoded, etc). M

Webhooks / Add webhook

Payload URL *

Webhook

- supports asynchronous & event-driven notifications
- Try Webhook.site.

Your unique URL

Your unique DNS name

Your unique email address

Proxy bidirectionally with Webhook site CLI

- Optionally, add a new GitHub Webhook in a repository to trigger send messages upon issue-related events.
 - Create a new issue and do sth there. New messages will appear at webhook.site.

https://webhook.site/be90c3aa-0079-454f-acee-3be4df1de409 P7 Open in new tab

be90c3aa-0079-454f-acee-3be4df1de409@emailhook.site ⋈ Open in mail client

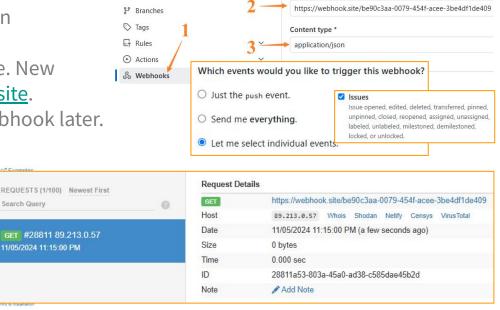
\$ whcli forward --token=be90c3aa-0079-454f-acee-3be4df1de409 --target=https://localhost

Remember to clean the GitHub Webhook later.

access this URL to simulate an event-driven message

Search Query

11/05/2024 11:15:00 PM



Reverse Proxy Flow

Summary

Reverse Proxy

- Reverse Proxy vs. Forward Proxy
- Benefits:
 - i. Load Balancing
 - ii. Access Control & Security
 - iii. Content Caching & Filtering
 - iv. Monitoring & Logging
 - v. SSL Encryption Offloading
- Nginx as a reverse proxy

Other API Architectures

- SOAP
- GraphQL
- WebSocket
- Webhook

That's all for this module!

