

# DISTRIBUTED AND CLOUD COMPUTING

LAB 9: MAP-REDUCE STREAMING & INTRO TO SPARK

# Plan for today

A more detailed look into YARN + YARN set up

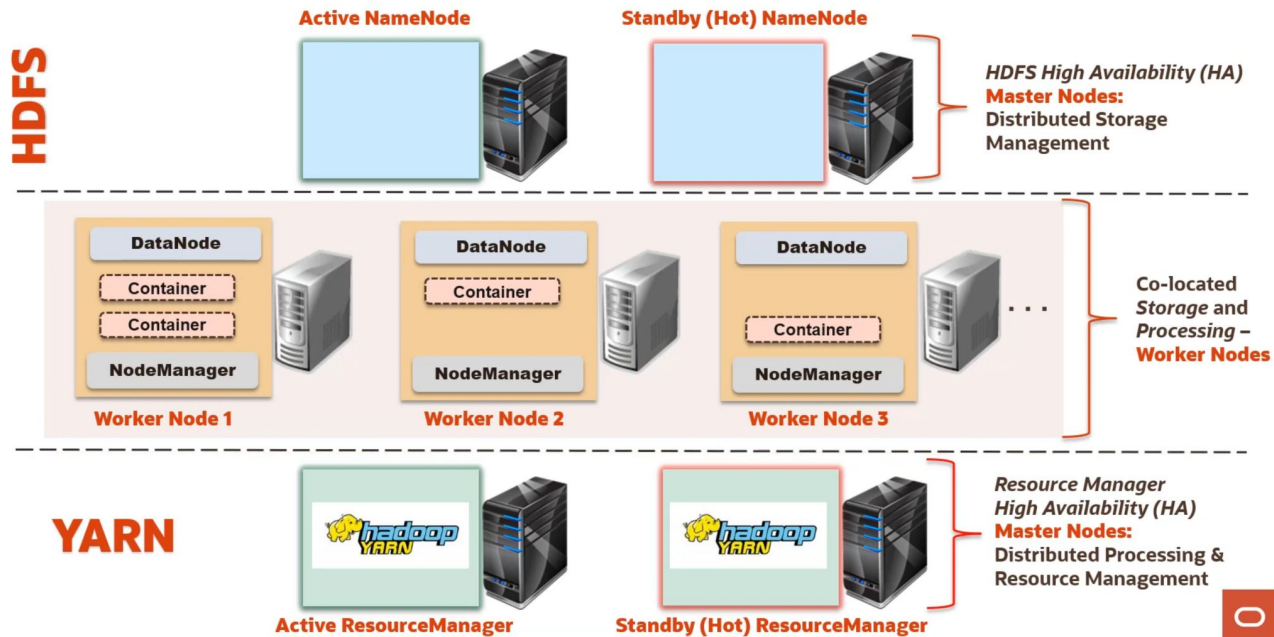
Mapreduce streaming & implementation of more complex mapreduce program

# YARN intuition

**YARN is to computation**  
**what HDFS is to data**  
**storage**

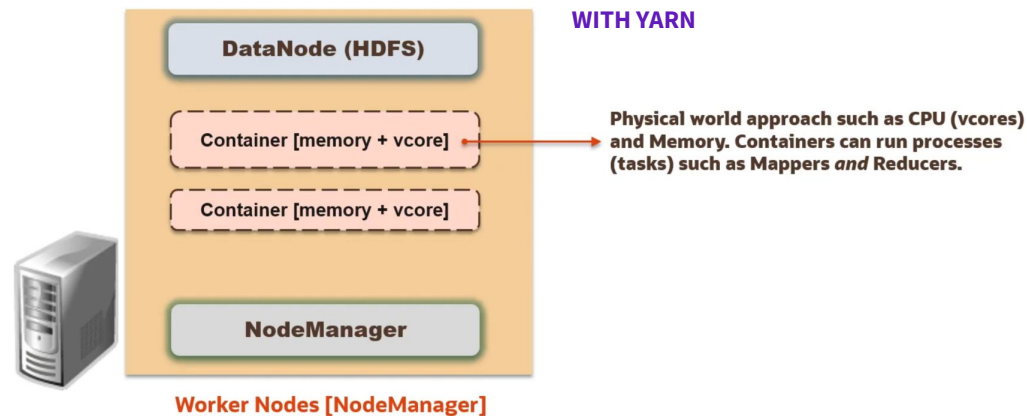
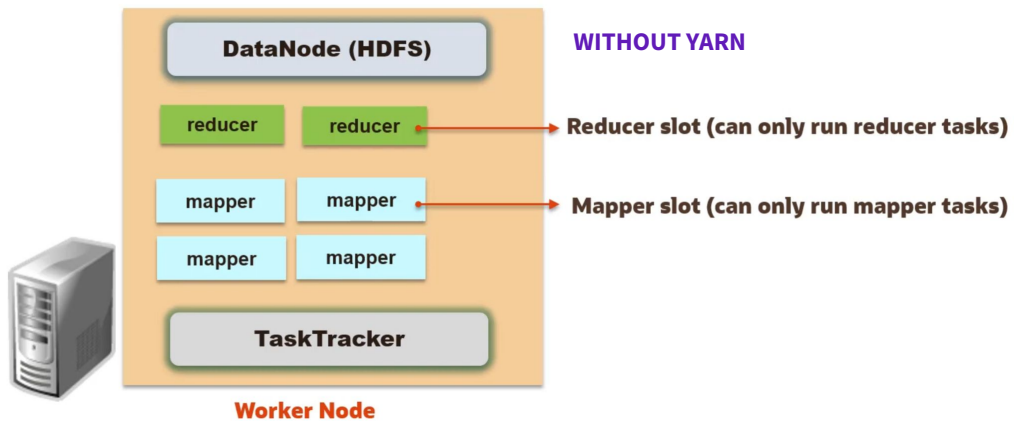
In HDFS a master node (NameNode) listens to client requests for file system operations and orchestrates the cluster to achieve them in a distributed fault-tolerant way

In **YARN** a master node (**ResourceManager**) listens to client requests for **computation (jobs)** and orchestrates the cluster to achieve them in a distributed fault-tolerant way




In HDFS the cluster computers run a DataNode daemon that listens to the NameNode  
In YARN the **SAME** cluster computers run a **NodeManager** daemon that listens to the ResourceManager

# Hadoop cluster computation with and without YARN

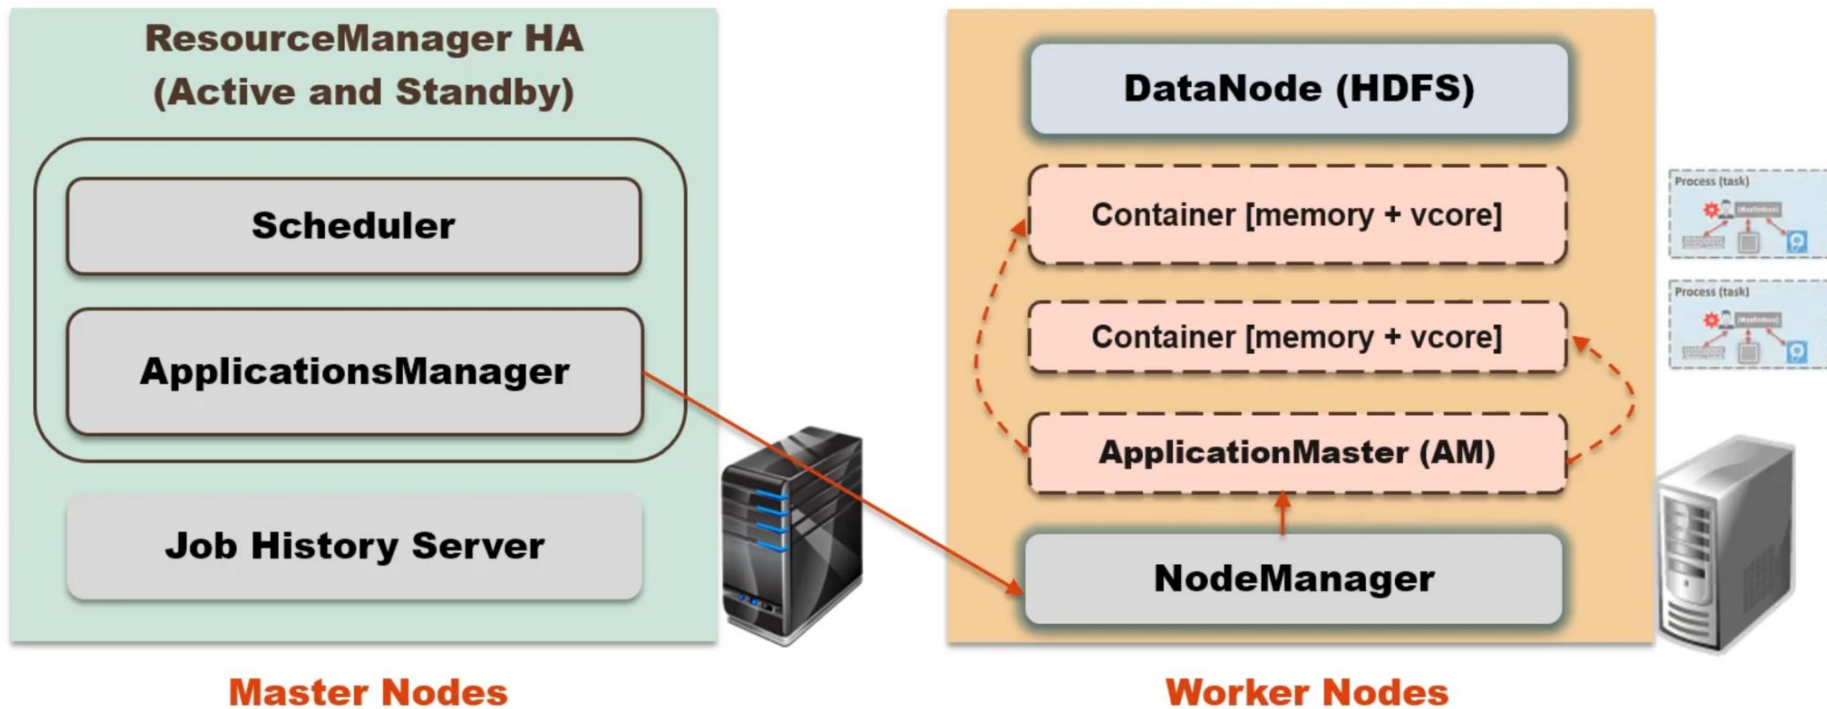


**YARN containers allow for execution of other applications besides Map Reduce**

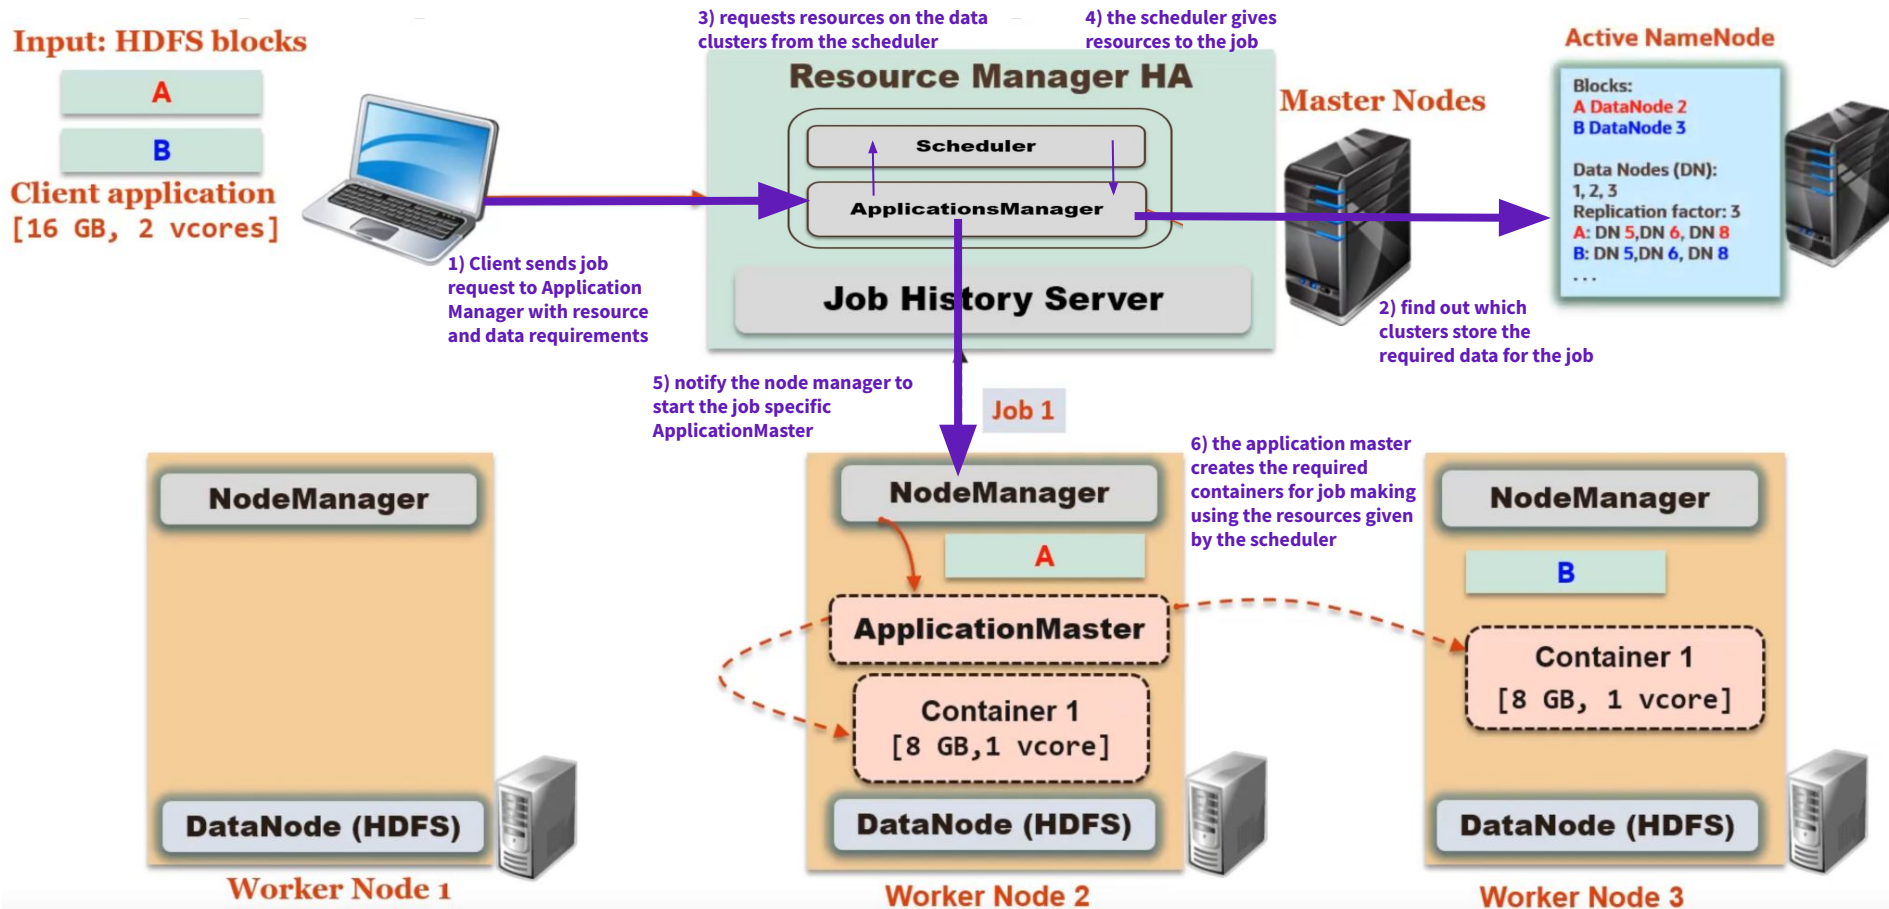
**Spark is the one we are interested in**



# YARN: Resource Manager and NodeManager



# YARN: Example of a YARN job execution



# YARN: Scheduling

- The YARN scheduler is in charge of determining **where** and **when** submitted jobs will be executed
- The following strategies are available for **when**

## First In, First out (FIFO) Scheduler

- Allocated resources to tasks in arrival order

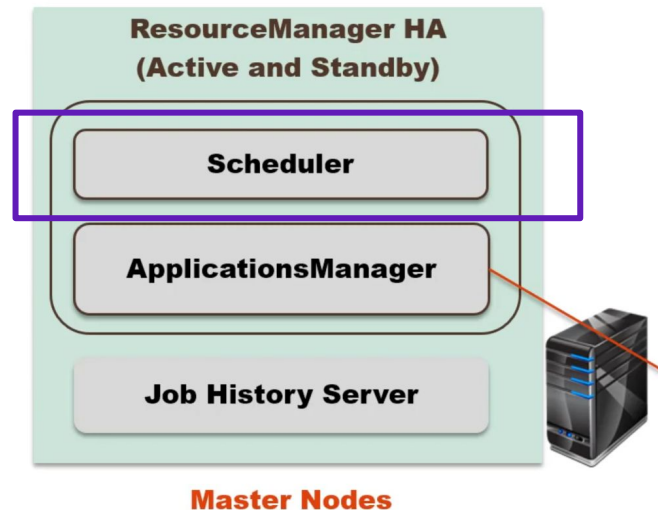
## Capacity Scheduler

- Tasks are assigned to priority queues
- Resources are assigned to tasks waiting in higher priority queues
- Within the same queue, FIFO is used

## Fair Scheduler

- Tasks are assigned to queues (NOT priority queues)
- The scheduler assigns resources to tasks in all queues FAIRLY

- The **where** is determined by **where** the data is stored in HDFS

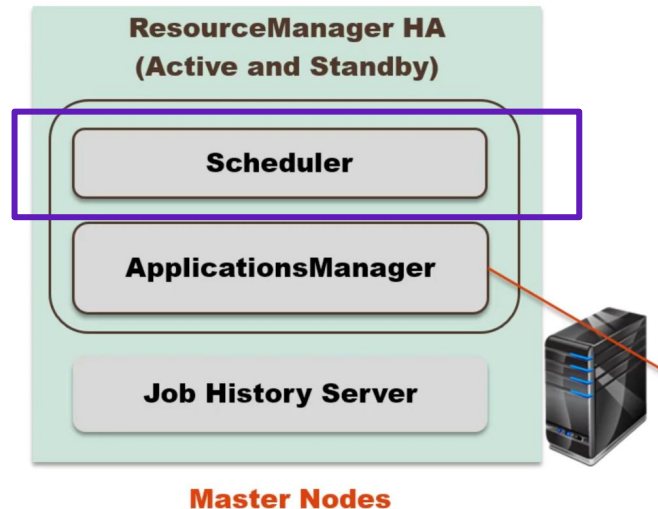


# YARN: Why is scheduling important

- Scheduling is extremely important
- An effective scheduling strategy can maximises the resource utilisation and minimises the task waiting time
- Bad resource utilisation → increased energy usage
  - 10 clusters running at 10% > 1 cluster at 100%
- However scheduling everything to one cluster will lead to massive waiting times

**Finding an optimal scheduling strategy for a given workload is an NP-Complete problem**

Fancy way of saying super hard





# TASK 1: Set up YARN in your local Hadoop installation

**STEP 1:** Start HDFS in pseudo-distributed mode

- Use this [link](#) to remember the commands that we used last time
- If you were not here last time follow the tutorial in last weeks

**STEP 2:** In the same link from STEP 1 navigate to the section:

“YARN on a Single Node”

- 1) Copy the configuration to the appropriate file xml configuration files
- 2) Execute \$ sbin/start-yarn.sh
  - a) You must be in the hadoop home directory
- 3) Navigate to localhost:8088
  - a) You should see something like the lower image

**STEP 3:** Execute a mapreduce operation and look around the YARN endpoint to see the execution details. You can do this by:

**A)** Modifying [this example](#) in standalone operation to work in the pseudo distributed mode by

- Creating the **HDFS://user/<YOU>** directory (if not there)
- Using hdfs commands to mimic **\$ cp etc/hadoop/\*.xml input** (place these files in **HDFS://user/<YOU>**)
- Execute the mapreduce command using the correct hadoop version and paths

**B)** Compiling and running the word count java program example [here](#)

- Be careful to modify the command with the correct paths etc..

## YARN on a Single Node

You can run a MapReduce job on YARN in a pseudo-distributed mode by setting a few parameters and running ResourceManager daemon and NodeManager daemon in addition.

The following instructions assume that 1. ~ 4. steps of the above instructions are already executed.

1. Configure parameters as follows:

etc/hadoop/mapred-site.xml:

```
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
  <property>
    <name>mapreduce.application.classpath</name>
    <value>${HADOOP_MAPRED_HOME}/share/hadoop/mapreduce/*:${HADOOP_MAPRED_HOME}/share/hadoop/mapreduce/lib/*</value>
  </property>
</configuration>
```

etc/hadoop/yarn-site.xml:

```
<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.env-whitelist</name>
    <value>JAVA_HOME,HADOOP_COMMON_HOME,HADOOP_HDFS_HOME,HADOOP_CONF_DIR,CLASSPATH_PREPEND_DISTCACHE,HADOOP_YARN_HOME,HADOOP_HOME,PATH,LANG,TZ,HADOOP_MAPRED_HOME</value>
  </property>
</configuration>
```

2. Start ResourceManager daemon and NodeManager daemon:

```
$ sbin/start-yarn.sh
```

3. Browse the web interface for the ResourceManager; by default it is available at:

ResourceManager - <http://localhost:8088/>



## All Applications

Cluster Metrics										Cluster Information														
Nodes Available	Nodes Pending	Nodes Running	Nodes Completed	Customers Training	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources	Nodes Resources					
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20					
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name					Cluster ID					Cluster Version					Cluster Status					Cluster Health				
Cluster Name																								

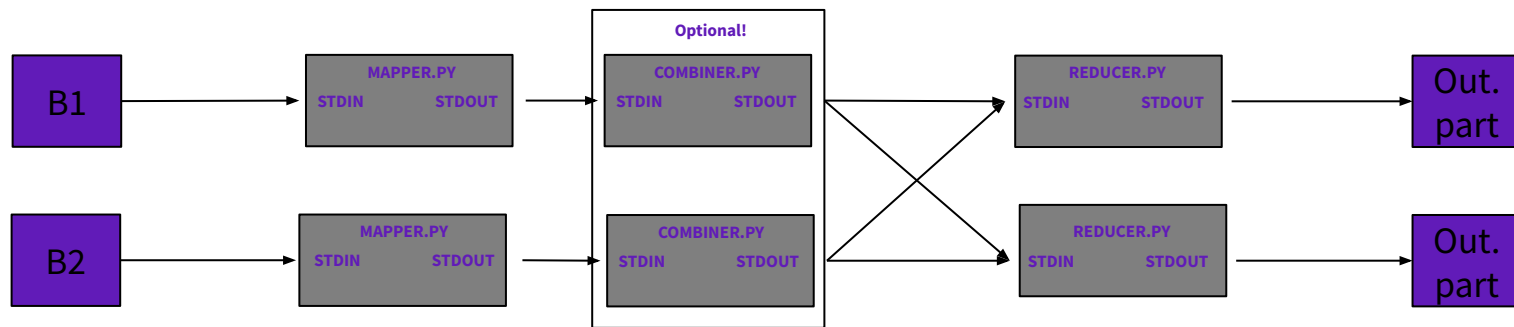
# MapReduce Streaming

- So far all the examples of MapReduce we have seen used Java binaries for the mapper and reducer processes
  - Using the pre-compiled examples that shipped with hadoop (grep)
  - Using mappers and reducers written jars compiled by us (word count)
- MapReduce streaming offers the ability to use **ANY** kind of executable for the mapper and the reducer process
  - These can be shell scripts, python scripts or any type of executable binary

**Gives the programmer much more freedom when writing mapreduce programs**

# MapReduce Streaming

- MapReduce streaming jobs utilise the **stdin** and **stdout** buffers
- Raw data are fed to the stdin of the mappers which output  $\langle k1, v1 \rangle$  to stdout
- The stdout of the mappers is sorted and fed to the stdin of the reducer scripts which output  $\langle k2, v2 \rangle$  to stdout
- The stdout of the reducers is written back to HDFS as parts
  - New Lines (`\n`) are used to separate the key value pairs
  - The first tab (`\t`) separates the key from the value
    - ***<key: text up to first tab, value: remaining text until new line>***
  - This can be customized through `-inputformat`



# Task 2: Mapreduce streaming with Python scripts

**TASK DESCRIPTION:** Write the **mapper.py** and **reducer.py** python scripts to be used by a mapreduce streaming job that searches 'commerce\_data\_cleaned.csv' for the most bought products by country (the format of the csv file can be seen below)

**STEP 0:** Start the NameNode, DataNodes and Yarn daemons and create **HDFS://user/<user\_name>** if not already there

**STEP 1:** Download the 'commerce\_data\_cleaned.csv' and streaming\_word\_count example from blackboard (dataset is from [here](#))

**STEP 2:** Move the .csv and .txt files to HDFS

**STEP 3:** Execute the providence mapper and reducer scripts with the input.txt to test that streaming works and its behaviour (also YARN)

**STEP 3:** Implement **mapper.py** and **reducer.py** (don't forget to include `#!/bin/python` at the first line in each file) for the tasks

**STEP 4:** Execute the mapreduce process on the 'commerce\_data\_cleaned.csv' and check out the results

Command for mapreduce streaming (assuming current working dir is the hadoop home directory)

```
$ /bin/mapred streaming -input <path_to_data> -output <path_to_output> -mapper <path_to_mapper.py> -file <path_to_mapper.py> -reducer <path_to_reducer.py> -file <path_to_reducer.py>
```

## CSV column labels

```
InvoiceNo,StockCode,Description,Quantity,InvoiceDate,UnitPrice,CustomerID,Country
536365,85123A,WHITE HANGING HEART T-LIGHT HOLDER,6,12/1/2010 8:26,2.55,17850,United Kingdom
536365,71053,WHITE METAL LANTERN,6,12/1/2010 8:26,3.39,17850,United Kingdom
```

## Snippet of the expected output

```
(base) george@DESKTOP-E24BUDU:~/DistSys24/hadoop/s
hdfs dfs -cat /user/george/output_commerce/*
Australia      ('MINI PAINT SET VINTAGE ', 2916)
Austria ('SET 12 KIDS COLOUR CHALK STICKS', 288)
Bahrain ('ICE CREAM SUNDAE LIP GLOSS', 96)
Belgium ('PACK OF 72 RETROSPOT CAKE CASES', 480)
```